

Temporal hierarchies in R

Nikolaos Kourentzes (nikolaos@kourentes.com)

Contents

1	Data and packages	1
2	Temporal hierarchies using the thief package	2
3	Manual implementation of THieF	3

1 Data and packages

We will model temporal hierarchies in two ways. First, we will use a package that gives us a somewhat limited control in the generation of the hierarchical forecasts. Second, we will use step by step code to do the combination manually, as it is a fairly simple process and gives us full control in generating any base forecasts we may want.

The relevant package is *thief*. For manual implementation we will make use of some supporting packages, namely: *MAPA* (this is in fact the 1st version of temporal hierarchies), *tsutils*, and *abind*.

Let us load the packages.

```
pckg <- c("thief", "MAPA", "tsutils", "abind")
for (i in 1:length(pckg)){
  if(!(pckg[i] %in% rownames(installed.packages()))){
    install.packages(pckg[i])
  }
  library(pckg[i], character.only = TRUE)
}

## Warning: package 'thief' was built under R version 4.0.5
## Loading required package: forecast
## Registered S3 method overwritten by 'quantmod':
##   method           from
##   as.zoo.data.frame zoo
## Loading required package: parallel
## Loading required package: RColorBrewer
## Loading required package: smooth
## Loading required package: greybox
## Package "greybox", v0.5.9 loaded.
## This is package "smooth", v2.5.6
## Warning: package 'tsutils' was built under R version 4.0.5
```

We will use the usual *AirPassenger* time series.


```
y <- AirPassengers
```

2 Temporal hierarchies using the thief package

Using the function `thief()` the whole process is automated:

```
frc1 <- thief(y)
plot(frc1)
```

Forecasts from THieF-ETS




By default it

uses ETS as the base forecast, we can change that using the argument `usemodel`. For instance:

```
frc2 <- thief(y,usemodel="arima")
plot(frc2)
```

Forecasts from THieF-ARIMA



Similarly, we can control the approximation of the covariance of the forecast errors that contribute to the estimation of the combination weights for the matrix G. This can be done with the argument `comb`. Look at the documentation of the function for more details `?thief`.

3 Manual implementation of THieF

First, we construct the summing matrix S, used to map the hierarchy and pick up from that the useful aggregation levels (we use only those that would not result in a fractional seasonality).

```
S <- tsutils::Sthief(y)      # Get the S matrix
ff <- frequency(y)          # Get sampling frequency of target series
AL <- ff/(1:ff)             # Calculate frequencies of various aggregation levels
AL <- AL[AL %% 1 == 0]       # And exclude those that would not be integer
k <- length(AL)             # Find how many are left
```

Next we create the temporally aggregated time series

```
Y <- MAPA::tsaggr(y,AL)[[1]]
```

We set the target forecast horizon, and calculate the corresponding forecast horizons for all temporal aggregation levels

```
hrz <- 16 # Target horizon
hAggr <- (ceiling(hrz/ff)*ff)/AL
hAggr
```

```
## [1] 2 4 6 8 12 24
```

Now we produce all the base forecasts, i.e., the independent forecasts for each level of the temporal hierarchy. For this example I am using ETS, but any forecasting method could be used. Note that this does not need to be the same across the various levels of the tmepral hierarchy.

```

frc <- mse <- list()
for (i in 1:k){
  yTemp <- Y[[i]]
  fit <- ets(yTemp)
  mse[[i]] <- fit$mse
  frcTemp <- forecast(fit,h=hAggr[i])$mean
  # Re-structure forecasts
  frc[[i]] <- matrix(frcTemp,ncol=hAggr[1]) # Organised as column per year
}

```

Next, we re/arrange the forecasts into a matrix. First comes the most aggregate, and the rest follow.

```

frcAll <- abind(frc,along=1)
frcAll

```

```

##          [,1]      [,2]
## [1,] 6117.6361 6596.1292
## [2,] 3007.6633 3250.5458
## [3,] 3129.1045 3371.9871
## [4,] 1833.0574 1941.5106
## [5,] 2393.3701 2532.2736
## [6,] 1898.4091 2006.5253
## [7,] 1355.3595 1459.7158
## [8,] 1593.3745 1713.7677
## [9,] 1887.3399 2027.3332
## [10,] 1383.1478 1483.8981
## [11,] 867.1471 929.9355
## [12,] 963.5666 1032.5139
## [13,] 1094.7423 1172.1632
## [14,] 1332.7008 1425.8650
## [15,] 1037.7271 1109.4450
## [16,] 881.9057 942.1689
## [17,] 441.8018 459.0139
## [18,] 434.1186 450.6333
## [19,] 496.6300 515.0797
## [20,] 483.2375 500.7700
## [21,] 483.9914 501.1423
## [22,] 551.0244 570.0974
## [23,] 613.1797 633.9130
## [24,] 609.3648 629.4938
## [25,] 530.5408 547.6630
## [26,] 463.0332 477.6340
## [27,] 402.7478 415.1573
## [28,] 451.9694 465.5780

```

Each column of frcAll corresponds to a period in the most temporally aggregate level, here the annual level. Rows contain the values for the forecasts at the various nodes of the hierarchy.

Now we estimate the G matrix that contains the combination weights. We will use here two easy approximations, the structural and the variance scaling.

```

# Structural:
W <- diag(1/rowSums(S))
Gstr <- solve(t(S) %*% W %*% S) %*% t(S) %*% W
# Variance:
mse <- unlist(mse)
W <- diag(1/mse[rep((1:k),rev(AL))])
Gvar <- solve(t(S) %*% W %*% S) %*% t(S) %*% W

```


Note that the calculation of G is the same in both cases. It is only the W that is different.

Now we reconcile (combine) the forecasts. For this example we use the Structural scaling, but we only need to replace G to obtain any other result.

```
# Create the bottom level forecasts
frcBRec <- Gstr %*% frcAll
frcFinal <- as.numeric(frcBRec)[1:hrz]
# We can also translate this into a time series object
frcFinal <- ts(frcFinal,frequency=frequency(y),start=end(y)[1] + deltat(y)*end(y)[2])
frcFinal

##          Jan      Feb      Mar      Apr      May      Jun      Jul      Aug
## 1961 440.8650 433.1818 491.3771 487.7213 509.7371 576.7701 646.6260 642.8111
## 1962 471.6625 463.2819 524.0035 520.9690
##          Sep      Oct      Nov      Dec
## 1961 547.0800 474.7431 411.4239 460.6455
## 1962

ts.plot(y,frcFinal,col=c("black","red"))
```



Or we can use the S-matrix to generate the forecasts for the complete hierarchy

```
frcARec <- S %*% frcBRec # Which has the same structure as frcAll
```

Happy forecasting!